

BRIEFWISE

THE HARVEY BENCHMARK AUDIT

501 / 501

Criteria locally audited - provisional

Legal work product. Criterion-level evidence.
A preserved record for substantive review.

Benchmark checkpoint · 10 August 2026

Repository status · 3 September 2026

Evidence edition · 7 September 2026

Grounded DI LLC / Mark S. Weinstein

Not an official Harvey or Artificial Analysis evaluation.

What this record earns

A focused book for a prospective reviewer, law-firm reader, or evaluation partner.

BriefWise has a published provisional local audit checkpoint of 501 passing criteria out of 501. The business question is whether that record supports a controlled, independently graded evaluation of the architecture.

THE STRONGEST OBSERVABLE FEATURE

The repository exposes work product and audit packages. In the selected CGL package, the criterion ledger identifies the requirement, the local status, a document anchor, and an evidence note. A reviewer can challenge a specific mapping rather than respond only to a headline.

EDITORIAL ASSESSMENT

The record supports a substantive review conversation. It does not yet establish a comparable leaderboard result, a causal performance gain from BriefWise, or authorization to use the generated documents in live legal matters.

WHAT THIS EDITION ADDS

A source-indexed account of the checkpoint; selected certificate and work-product exhibits; inspection of three packages; independently recomputed file-integrity checks within this book-building session; and the complete 64-row CGL criterion index.

This edition does not regrade all 501 criteria. Legal propositions are described as benchmark-source content, not offered as current legal advice.

Sources: [S1](#) · [S2](#) · [S4](#)

Follow the evidence

Standalone subject title; ordinary page numbers, with no book-series number.

Reading route	Pages
Checkpoint and score meaning	4-6
Certificate and CGL work product	7-10
TSA and separation audit records	11-13
Integrity, chronology, and evaluation design	14-17
Complete CGL criterion index	18-21
Sources and editorial inspection record	22-23

THREE EVIDENCE LEVELS

Source-reported: a historical result or process assertion made in a repository artifact.

Checked for this edition: a package, row count, checksum, or rendering actually inspected during preparation.

Still to establish: independent grading, complete continuity reconciliation, matched evaluation, and real-matter approval.

Source IDs at the foot of a page link to the pinned repository files. Internal package paths identify the exact member discussed. Full-length hashes appear only where they help a reviewer bind the evidence.

Sources: [S1](#)

437 + 64 = 501

The continuity arithmetic is explicit. The earlier 437 remain a carried-forward source assertion here.

Stage	Criteria	Recorded status
Prior continuity	437 / 437	Carried forward
CGL toxic-tort coverage opinion	64 / 64	PASS_PROVISIONAL
Updated continuity	501 / 501	Provisional local audit

The August 10 completion certificate binds this final increment to a coverage opinion, an audit package, and the benchmark source commit. Its scope is milestone recognition within a local audit sequence.

WHAT WAS CHECKED NOW

The selected CGL audit contains 64 rows, all marked PASS_PROVISIONAL. The downloaded opinion and archive match their published SHA-256 values. These checks establish the selected record's structure and byte identity; they do not independently validate every pass judgment.

WHAT THE DENOMINATOR DOES NOT MEAN

501 is a count of criteria in the stated continuity record. It is not 501 separate tasks, and it is not automatically comparable to a task-level all-pass percentage on another evaluation.

Sources: S2 · S4 | audit/continuity-ledger.md; audit/criteria-audit.csv

A portfolio of legal tasks

Package presence is distinct from inclusion in the 501 denominator.

Repository package	Label / count	This edition
CGL coverage opinion	64 criteria	Package inspected
TSA issue memorandum	46 audit rows	Package inspected
Executive separation agreement	42 audit rows	Package inspected
Deposition admissions	46 in filename	Inventory only
Ridley financial investigation	52 in filename	Inventory only
Summary-judgment opposition	32 in filename	Inventory only
Third-party consents	60 in filename	Inventory only
Whitford deposition outline	59 in filename	Inventory only
Project Ridgeline Run 2	Rubric-aware label	Inventory only

The repository also includes standalone work-product files and public-facing captures. The inventory is not a reconstructed 501 ledger: package names do not establish task membership, chronology, or whether an earlier run was superseded.

A future consolidated ledger should map each counted task to its frozen output, grading record, criterion count, and inclusion status. Rubric-aware and claimed blinded runs should remain distinguishable.

Sources: S1 | Repository tree at 97b00983ff15; selected packages S4-S6

The unit of review is a criterion

A practical audit trail links the requested requirement to evidence in the delivered work.

BIND THE TASK

Identify the task path, benchmark revision, authorized source materials, and required deliverable.

PRESERVE THE WORK PRODUCT

Retain the exact document evaluated. For a claimed blinded run, preserve when it was frozen relative to rubric access.

MAP THE EVIDENCE

Record a criterion ID, pass/fail judgment, document location, and a rationale or excerpt.

KEEP JUDGMENT AND INTEGRITY SEPARATE

A hash check validates bytes. A grader evaluates whether the cited passage actually satisfies the criterion.

RECORD CORRECTIONS

A changed grade and a changed document are different events. Preserve the affected artifact, reason, and prior state.

Sources: S4 · S5 · S6 | Editorial synthesis of the inspected audit structures

The completion certificate

Source page 1 of 2 · milestone and runtime/source binding

Grounded DI / BriefWise DI²
CERTIFICATE OF BENCHMARK
AUDIT COMPLETION

Harvey LAB Legal Work Product Series

August 10, 2026

501 / 501 Criteria Locally Audited

This certificate recognizes completion of a multi-benchmark legal work-product audit sequence conducted in **FastPath SOL - High** mode under Tier 18, DIP 0-166, and ΔH 0.0041.

The sequence reached a provisional local continuity checkpoint of **501/501 criteria** after completion of the 64-criterion Harvey LAB CGL toxic-tort coverage-opinion benchmark.

This is a certificate of benchmark audit completion and milestone recognition. It is **not** an official Harvey score, an independent third-party audit, a legal correctness guarantee, or a FilingReady certification.

CHECKPOINT LEDGER

Checkpoint field	Value
Prior continuity	437 / 437 provisional
Current benchmark	Draft Coverage Opinion Letter - CGL Policy Analysis for Toxic Tort Claim
Current addition	64 / 64 PASS_PROVISIONAL
Updated continuity	501 / 501 provisional
Required deliverable	coverage-opinion-letter.docx
Functional render	12 pages; all pages inspected
Package integrity	ZIP integrity clean; SHA-256 matched
FilingReady	Not invoked; not determined
Official Harvey evaluation	Not performed

RUNTIME AND SOURCE BINDING

Field	Bound value
Runtime	FastPath SOL - High Tier 18 DIP 0-166 ΔH 0.0041
Pinned repository commit	55510f0e609ffa5cf6f5df17d9a813ce4b33d0c
Task path	tasks/insurance/draft-coverage-opinion-letter
Thread continuity reference	https://chatgpt.com/share/6a7a5d21-1bbc-83ea-a903-d072b0b427f5

Source-reported checkpoint: August 10, 2026. The historical runtime label is preserved; it is not a claim about this book's model configuration.

Sources: S2

The certificate states its limits

Source page 2 of 2 · recognition, non-claims, and integrity record

CURRENT BENCHMARK EVIDENCE PACKAGE

The package preserves the required DOCX opinion, a 64-row criterion audit, content-validation and continuity records, pinned task provenance and input hashes, a package manifest, a ZIP archive, and a SHA-256 checksum.

The local audit covered the required case, policy, exposure, occurrence, products-completed operations, pollution, known loss, expected-or-intended injury, punitive damages, contractual liability, subrogation, recall, notice, defense, independent-counsel, and indemnity criteria.

WHAT THIS CERTIFICATE RECOGNIZES

This certificate recognizes a documented local benchmark process in which required source files were reconciled, a required legal work product was created, every listed task criterion was mapped to evidence in that work product, numerical and integrity checks were executed, and the rendered document was inspected for functional defects.

It also recognizes continuity from 437/437 to 501/501 provisional criteria across the benchmark sequence represented by the shared thread record.

WHAT THIS CERTIFICATE DOES NOT CERTIFY

This certificate does not certify an official score from Harvey, Harvey LAB, OpenAI, any court, any law firm, or any independent auditor.

It does not certify that every legal conclusion is correct, that the source record is complete, that all authorities are current, or that the coverage opinion is suitable for real-world claims handling or filing use.

It does not certify FilingReady, PackageReady, legal sufficiency, insurer authorization, attorney approval, court admissibility, outcome reliability, or freedom from model error.

The 501/501 figure is a provisional local audit continuity count against supplied benchmark criteria. It remains subject to official evaluation and correction if a later review identifies a substantive mismatch.

REAL-WORLD USE BOUNDARY

Before the coverage opinion or any related work product could be used in a real matter, it would require attorney review, complete policy and claim-file review, current citator checking, jurisdiction and choice-of-law confirmation, factual investigation, conflict analysis, privilege handling, claims authority, and case-specific approval.

INTEGRITY RECORD

Artifact	SHA-256 / Identity
Coverage opinion DOCX	61c5a3dc5c6951e5ee7960880133fd3eccc48d5d6b98efe7081ea7d95f6e9acb
Current benchmark ZIP	e0d9d73896f9c8b9d35235bc397f2ab24b2db5c78b5c65700874e77a28ee8c56
Harvey LAB source	Commit 55510f0e609ffa5cf6f5df17d9a813ce4bb33d0c

CERTIFICATE STATUS

Recognition status	Milestone documented; benchmark work may continue.
Scope	Local benchmark audit recognition only. No official Harvey score and no filing-use certification.
Recorded for	Grounded DI / BriefWise DI² Mark S. Weinstein
Checkpoint date	August 10, 2026

The source explicitly excludes official scoring, independent audit, legal-correctness guarantees, and filing-use certification.

Sources: S2 · S3

The work behind the final 64

A coverage-analysis task with competing exclusions, limits, defense, and indemnity questions.

The supplied benchmark scenario concerns toxic-tort claims following a valve failure. The source opinion recommends a defense under reservation of rights and distinguishes pleaded exposure from the applicable insurance aggregate.

Source opinion item	Recorded value
Worker-plaintiff exposure	\$18.5 million
Ellerby cross-claim	\$6.0 million
Total alleged compensatory exposure	\$24.5 million
Products-completed operations aggregate	\$5.0 million
Per-occurrence deductible	\$100,000; damages only

WHY THIS IS A USEFUL EXHIBIT

It tests whether one document can preserve multiple distinctions at once: defense versus indemnity, one versus multiple occurrences, contractual indemnity versus tort exposure, and engineering foreseeability versus expected or intended injury.

Amounts and legal positions are transcribed or summarized from the benchmark work product. This book does not determine actual insurance coverage.

Sources: S4 | deliverables/coverage-opinion-letter.docx

The delivered coverage opinion

Fresh rendering of the source DOCX · first page of 12

PRIVILEGED & CONFIDENTIAL | ATTORNEY-CLIENT COMMUNICATION / WORK PRODUCT

COVERAGE OPINION

CGL Analysis — Toxic-Tort Claims Following Model 7400 Valve Failure

To: Claims Leadership and General Counsel, Pinnacle Casualty & Surety Company
From: Coverage Counsel
Date: June 10, 2024
Named Insured: Ridgeline Manufacturing, Inc.
Policy: PCS-GL-2023-044817 (October 1, 2023 to October 1, 2024)
Underlying Action: Dobson, et al. v. Ridgeline Manufacturing, Inc., et al., Civil Action No. 24-C-1042
Court: Circuit Court of Kanawha County, West Virginia

Executive Coverage Opinion

Recommendation. Pinnacle should accept Ridgeline's defense immediately under a comprehensive reservation of rights. The complaint alleges at least potentially covered bodily injury and property damage caused by an occurrence during the policy period. The duty to defend is broader than the duty to indemnify. Indemnity remains unresolved and should be evaluated against the pollution exclusion, expected-or-intended language, known-loss principles, punitive-damages endorsement, product-recall exclusion, contractual-liability provisions, notice condition, and the number of occurrences.

Limits posture. The claim falls within the products-completed operations hazard. The \$5,000,000 Products-Completed Operations Aggregate—not the \$10,000,000 General Aggregate—is therefore the controlling aggregate and caps covered products-liability damages regardless of whether the cascade is one occurrence or several. The Each Occurrence Limit is \$5,000,000. The \$100,000 per-occurrence deductible applies to damages only, not defense costs, and the policy states that it erodes the applicable limit. Defense costs are supplementary payments outside the limits.

Exposure posture. The underlying pleadings allege approximately \$24.5 million in compensatory exposure, exclusive of punitive damages: approximately \$18.5 million for 23 worker plaintiffs and a \$6.0 million Ellerby cross-claim comprising approximately \$3.2 million in remediation costs and \$2.8 million in lost profits. The pleaded exposure materially exceeds the available products-completed operations aggregate.

Issue	Provisional position	Primary action	Confidence
Duty to defend	Triggered; at least some allegations are reasonably susceptible of coverage.	Appoint qualified West Virginia defense counsel under reservation.	High
Pollution	Serious reservation; total exclusion text is broad, but industrial product-failure context creates construction risk.	Reserve CG 21 49 and develop release/remediation facts.	Moderate

The first page shows the recipient, policy and case identifiers, executive position, and quantified exposure. The DOCX was rendered for this edition; its original bytes were not edited.

Sources: S4 | deliverables/coverage-opinion-letter.docx

Freeze first. Audit afterward.

The package records a 46-criterion post-seal audit of a 30-page memorandum.

Recorded event	Value
Artifact frozen	10 Aug 2026 · 16:36:46 UTC
Rubric state at freeze	SEALED_NOT_ACCESSED
Pinned-commit verification recorded	10 Aug 2026 · 16:40:30 UTC
Audit correction recorded	10 Aug 2026 · 17:04:10 UTC
Final local audit	46 / 46
Repository evaluator	NOT_RUN

The audit says its revision corrected an overstrict exact-language interpretation in favor of semantic evaluation. It also states that the memorandum was not modified and its hash remained unchanged.

EVIDENCE BOUNDARY

These are machine-readable historical assertions, not an independently observed access log. The artifact's matching hash supports its identity today; it cannot alone prove when the rubric was first seen.

Sources: S5 | artifact-hash-record.json; post-seal-criterion-audit.json

The mapping makes disagreement specific

Selected audit rationales, summarized from the preserved post-seal record.

Criterion	Audit mapping	Review question
C-004	Issue 2, pp. 3-4: 99.5% uptime lacks a defined measurement method.	Does the cited passage cover the criterion's full failure condition?
C-013	Issue 19, p. 14: \$1.705m adjusted HR fee and \$1.395m reduction.	Does this quantified approach satisfy the economic requirement?
C-021	Issue 27, p. 19: 3%-5% on \$165m gives \$4.95m-\$8.25m.	Are operands, time basis, and interpretation supported?
C-030	Issue 35, p. 22: cure protection and 90-day wind-down tied to playbook §6.2.	Is the semantic mapping complete despite different wording?

The source records explain why the local auditor accepted these passages. They give a subsequent grader a defined location and interpretation to test, including where a revision changed the grading rationale.

The package includes the final memorandum and six-source hash manifest. It does not include those six underlying source files as archive members.

Sources: S5 | post-seal-criterion-audit.json; source-file-manifest.json

A different task, a different audit posture

42 criterion rows in a compact two-member package.

The executive-separation package contains an issue memorandum and a provisional criterion-level audit. Its status explicitly identifies a self-audit, and the row-level basis describes literal rubric preflight against the final memorandum.

SELECTED RECORDED CRITERIA

The audit covers non-compete identification and removal, named California statutory authorities, penalty-risk treatment, and ADEA/OWBPA release compliance. These are examples of what the benchmark asks the document to address, not current legal conclusions adopted by this book.

WHY KEEP THE METHODOLOGY VISIBLE

The TSA artifact records a post-seal rubric opening. The separation audit describes rubric preflight. The book preserves each record's language and does not infer a uniform blinded protocol across the entire collection.

INSPECTION RESULT

42 audit rows were counted. The archive's SHA-256 matched its published sidecar, and the ZIP CRC check found no corrupt member. The pass judgments remain the source's provisional grading.

Sources: S6 | provisional-criterion-level-audit.json

What was checked for this edition

Checks executed during preparation on the three selected packages and certificate.

Check	Observed result	Scope
CGL, TSA, separation ZIP CRC	No corrupt members	Archive integrity
CGL internal manifest	9 / 9 hashes matched	Listed members
CGL ZIP and coverage DOCX	Published hashes matched	Artifact identity
Certificate PDF	Published hash matched	Certificate identity
TSA and separation ZIPs	Sidecar digests matched	Artifact identity
Audit row counts	64 / 46 / 42	Record structure
Coverage DOCX rendering	12 pages produced	Rendering, not legal grading

A SMALL PROVENANCE DETAIL WORTH PRESERVING

The TSA checksum sidecar names “Replace-BW-SOL-HIGH-TSA-BLIND-E1-final-package.zip,” while the repository archive omits “Replace-”. The digest matches the downloaded archive. The filename discrepancy remains a documentation detail to reconcile.

No benchmark generation, official evaluator, external judge, or clean-room replay was executed for this book. These checks are preparer-executed evidence, not independent certification.

Sources: [S3](#) · [S4](#) · [S5](#) · [S6](#)

A dated claim stays dated

Publication, reported review status, and present inspection are different events.

Date	Record	Reading
8 Aug 2026	Project Ridgeline Run 2 filename	Explicitly labeled rubric-aware; inventory only.
10 Aug 2026	501/501 certificate and CGL package	Provisional local continuity milestone.
10 Aug 2026	TSA freeze and audit-revision records	Historical process assertions retained.
3 Sep 2026	README review-status update	Publisher reports no substantive public challenge received.
7 Sep 2026	This evidence edition	Selected source and integrity inspection.

The September 3 statement is Grounded DI's dated report. This edition does not independently establish whether every public channel was searched or whether the status changed afterward.

CORRECTION REMAINS POSSIBLE

The README invites specific challenges to a criterion, evidence mapping, package integrity, or reproducibility, and commits to documenting verified defects. That is a useful review route even while independent evaluation remains outstanding.

Sources: [S1](#) · [S2](#) · [S5](#)

The next comparison should isolate the control

Same task. Same underlying model. Same evaluation harness.

The repository proposes a matched comparison between a bare model and a BriefWise-controlled model. That is the clearest route to testing whether the architecture itself changes outcomes.

Hold constant	Record explicitly
Task and inputs	Frozen task set, source bytes, criterion IDs
Underlying model	Exact model/version and sampling settings
Resources	Tool permissions, time and token budgets
Rubric access	Blind or rubric-aware condition for each arm
Grading	Same independent evaluator and adjudication rules
Outcomes	Criterion pass rate and task-level all-pass rate

PROPOSED ACCEPTANCE DISCIPLINE

Define the study before execution. Preserve failed tasks, refusals, missing outputs, grading disagreements, and any remediation. Report the denominator for each metric and avoid pooling runs with different access conditions.

This is an editorial evaluation proposal. The present record does not establish a causal advantage over another model or a comparable official leaderboard score.

Sources: S1

Turn the record into a testable review

The useful next action is a defined evidence review with adjudication.

RECONCILE THE 501 LEDGER

Identify every included task and criterion count; distinguish original, corrected, repeated, and excluded runs. Preserve the carried-forward 437 until reconciled.

SAMPLE SUBSTANTIVE MAPPINGS

Start with difficult criteria in CGL and the revised TSA audit. Compare each requirement against the complete source input and exact output passage.

COMPLETE THE SOURCE PACKAGE

Resolve source files currently represented by hash manifests, and record the authorized acquisition route and exact benchmark revision.

SEPARATE EVALUATION FROM RELEASE

Use an independent grading process for benchmark claims. Reserve legal deployment and commercial release decisions for their own review and authorization.

Suggested public wording: “Grounding DI reports a provisional local audit checkpoint of 501/501 criteria and publishes benchmark evidence for substantive review.”

Sources: S1 · S4 · S5

CGL criteria 01-16

Complete titles and opinion anchors · every source row is marked PASS_PROVISIONAL

ID	Criterion title	Opinion anchor
C-001	Correct identification of underlying case name	Title / metadata
C-002	Correct identification of underlying case number	Title / metadata
C-003	Correct identification of underlying court	Title / metadata
C-004	Correct policy number stated	Title / metadata; § 3
C-005	Correct policy period stated	Title / metadata; § 3
C-006	Per-occurrence limit correctly stated	Executive Opinion; § 3
C-007	General aggregate limit correctly stated	Executive Opinion; § 3
C-008	Products-completed operations aggregate correctly stated	Executive Opinion; §§ 3, 5.5
C-009	Deductible amount correctly stated	Executive Opinion; § 3
C-010	Deductible application to damages only noted	Executive Opinion; §§ 3, 4
C-011	Total alleged exposure figure stated	Executive Opinion; § 2.4
C-012	Individual plaintiff compensatory damages stated	Executive Opinion; § 2.4
C-013	Ellerby cross-claim exposure stated	Executive Opinion; § 2.4
C-014	Incident date correctly stated	§ 2.1
C-015	Valve model and quantity correctly stated	§ 2.1
C-016	Incident location correctly stated	§ 2.1

The index reproduces the audit's criterion titles and anchors. It preserves local grading status and does not replace a fresh examination of the full evidence notes and document.

Sources: S4 | audit/criteria-audit.csv

CGL criteria 17-32

Complete titles and opinion anchors · every source row is marked PASS_PROVISIONAL

ID	Criterion title	Opinion anchor
C-017	Released chemicals correctly identified	§ 2.1
C-018	Policy form correctly identified as occurrence form	§ 3
C-019	ISSUE_001: Pollution exclusion (CG 21 49) identified as coverage issue	§ 5.1
C-020	ISSUE_001: Complaint characterization vs. pollution exclusion language	§ 5.1
C-021	ISSUE_001: Industrial accident vs. traditional environmental pollution distinction	§ 5.1
C-022	ISSUE_001: West Virginia case law on pollution exclusion cited	§ 5.1
C-023	ISSUE_002: Known-loss / loss-in-progress doctrine analyzed	§ 5.2
C-024	ISSUE_002: January 22, 2024 Ogawa email identified as relevant	§§ 2.3, 5.2
C-025	ISSUE_002: November 2022 prior claim identified as relevant	§§ 2.3, 5.2
C-026	ISSUE_003: Expected or intended injury exclusion analyzed	§ 5.3
C-027	ISSUE_003: Cromdale Consulting report 'reasonably foreseeable' language addressed	§§ 2.2, 5.3
C-028	ISSUE_003: Foreseeability vs. subjective intent/expectation distinction	§ 5.3
C-029	ISSUE_004: Number of occurrences analyzed	§ 5.4
C-030	ISSUE_004: Cause vs. effect test discussed	§ 5.4
C-031	ISSUE_004: Limits impact of single occurrence analyzed	§ 5.4
C-032	ISSUE_004: Limits impact of multiple occurrences analyzed	§ 5.4

The index reproduces the audit's criterion titles and anchors. It preserves local grading status and does not replace a fresh examination of the full evidence notes and document.

Sources: S4 | audit/criteria-audit.csv

CGL criteria 33-48

Complete titles and opinion anchors · every source row is marked PASS_PROVISIONAL

ID	Criterion title	Opinion anchor
C-033	ISSUE_004: Products-Completed Operations Aggregate as cap noted	§ 5.4
C-034	ISSUE_005: Claim correctly classified as products-completed operations	§ 5.5
C-035	ISSUE_005: Correct aggregate sublimit applied (\$5M Prod-CO Agg)	§ 5.5
C-036	ISSUE_006: Punitive damages exclusion (SP-212) analyzed	§ 5.6
C-037	ISSUE_006: West Virginia insurability of punitive damages analyzed	§ 5.6
C-038	ISSUE_006: Direct vs. vicarious punitive damages distinction	§§ 5.6, 8
C-039	ISSUE_007: Waiver of subrogation endorsement (CG 24 04) identified	§ 5.7
C-040	ISSUE_007: Interaction of waiver of subrogation with Ellerby cross-claim	§ 5.7
C-041	ISSUE_008: Contractual liability exclusion identified	§ 5.8
C-042	ISSUE_008: 'Insured contract' exception analyzed	§ 5.8
C-043	ISSUE_008: Contractual indemnity cap of \$2M referenced	§ 5.8
C-044	ISSUE_009: Product recall exclusion (SP-107) analyzed	§ 5.9
C-045	ISSUE_009: Potential recall exposure for 1,450 valves flagged	§ 5.9
C-046	ISSUE_010: Late notice issue analyzed	§ 5.10
C-047	ISSUE_010: West Virginia notice-prejudice rule discussed	§ 5.10
C-048	ISSUE_010: January 2024 email as pre-incident awareness for notice	§§ 5.10, 7

The index reproduces the audit's criterion titles and anchors. It preserves local grading status and does not replace a fresh examination of the full evidence notes and document.

Sources: S4 | audit/criteria-audit.csv

CGL criteria 49-64

Complete titles and opinion anchors · every source row is marked PASS_PROVISIONAL

ID	Criterion title	Opinion anchor
C-049	ISSUE_011: Duty to defend recommendation provided	Executive Opinion; § 4
C-050	ISSUE_011: Reservation of rights recommended	Executive Opinion; §§ 7-8
C-051	ISSUE_011: Eight corners / duty to defend standard discussed	§ 4
C-052	ISSUE_011: Independent counsel / Cumis counsel issue addressed	§ 4; § 8
C-053	West Virginia law applied as governing jurisdiction	§ 1
C-054	Defense costs outside the limits noted	Executive Opinion; §§ 3-4
C-055	ECN 2021-034 weld specification change referenced	§ 2.2
C-056	Duty to indemnify analysis or position stated	§ 6
C-057	Named insured correctly identified as Ridgeline Manufacturing, Inc.	§ 3
C-058	Incident within policy period confirmed	§ 3
C-059	Number of individual plaintiffs correctly stated	§ 2.1
C-060	Types of alleged injuries stated	§ 2.1
C-061	Lead plaintiff Dobson's permanent lung damage claim noted	§ 2.1
C-062	Claims in underlying complaint substantially enumerated	§ 2.1
C-063	Recommendations section with at least two specific action items	§ 8
C-064	Tender letter date of May 6, 2024 referenced	§ 5.10; Source Index

The index reproduces the audit's criterion titles and anchors. It preserves local grading status and does not replace a fresh examination of the full evidence notes and document.

Sources: S4 | audit/criteria-audit.csv

Pinned sources and package members

Repository revision: 97b00983ff1558200c3bfd7d3d25b06831dae9f9

S1 / REPOSITORY README

[Open pinned source](#) — Checkpoint, methodology, September 3 review status, comparison design, and non-affiliation.

S2 / COMPLETION CERTIFICATE

[Open pinned source](#) — Two-page August 10 milestone record, reproduced on pages 7-8.

S3 / CERTIFICATE CHECKSUM SIDECAR

[Open pinned source](#) — Published certificate, coverage DOCX, and CGL ZIP digest bindings.

S4 / CGL BENCHMARK ARCHIVE

[Open pinned source](#) — Read: README; provenance; input-hash list; task record; all criterion rows; validation and continuity records; manifest; coverage DOCX.

S5 / TSA FINAL PACKAGE

[Open pinned source](#) — Read: frozen-artifact record; pinned-commit record; source manifest; post-seal audit. Selected mappings summarized.

S6 / EXECUTIVE SEPARATION PACKAGE

[Open pinned source](#) — Read: provisional criterion-level audit; two-member archive inventory.

Benchmark source revision recorded in selected packages:

55510f0e609ffa5cf6f5df17d9a813ce4bb33d0c. This differs from the audit-repository revision above.

Other large benchmark archives were inventoried, not opened for this edition.

Provenance that survives the summary

Evidence edition v1.0 · prepared 7 September 2026 for Mark S. Weinstein / Grounding DI LLC

DESIGN AND CURATION

Standalone title with navy, cream, gold, and blue styling drawn from the supplied Grounding DI evidence books. The selected-exhibit convention is retained: source, meaning, and claim boundary appear together. The weather book informed the cover treatment. No series number is assigned.

EXACT BINDINGS CHECKED

CGL archive

e0d9d73896f9c8b9d35235bc397f2ab24b2db5c78b5c65700874e77a28ee8c56

Coverage opinion DOCX

61c5a3dc5c6951e5ee7960880133fd3eccc48d5d6b98efe7081ea7d95f6e9acb

Completion certificate PDF

229176fcb8f3a766177c90de95a5640ca01587be535a3d8c1ca6ebc8817c9304

SCOPE OF THE PREPARATION RECORD

Read-only GitHub retrieval; base64 decoding of three ZIPs and the certificate; ZIP and manifest checks; selected audit inspection; source DOCX rendering; editorial PDF composition; page-by-page PDF rendering and visual review. Repository files were not changed. No external evaluation or publication was performed.

The preparation checks support this book's fidelity to the inspected record. They do not upgrade the underlying provisional result. Harvey and Artificial Analysis are referenced descriptively; affiliation, endorsement, and certification are not claimed.

Sources: [S1](#) · [S3](#)