

THE EVIDENCE ARC

From deterministic ideas to public artifacts



A plain-language record of the posts, diagrams, demonstrations, certificates, and careful limits that shaped Grounded DI public work from May 2025 through August 2026.

16 selected exhibits | monthly and weekly public repositories | curation freeze: 11 August 2026

The source record is preserved. The explanation is intentionally simple.

Precision does not guess. It traces.

CURATION NOTE

This book is a story about what became visible.

The social repositories contain far more material than one book can carry. This selection keeps the impressive moments, but explains them in ordinary language.

SOURCE PRESERVATION

Each exhibit names the public repository, file, and source commit used for this curation pass.

THE SIMPLE VERSION

The pages explain what happened before they explain why it matters. Technical terms are spelled out when they first appear.

CLAIM BOUNDARIES

A screenshot can show what was posted. A hash can show that an artifact stayed the same. Neither one automatically proves every larger claim.

WHAT THIS IS NOT

This is not a patent opinion, an independent benchmark evaluation, a legal certification, or a claim that every public statement has been externally verified.

The editorial rule is simple:

Use the strongest sentence the source actually earns.

THE STORY IN ONE PAGE

The system kept getting easier to explain.

The public record moves from an idea, to an architecture, to modules, to controlled outputs, and finally to evidence that another person can inspect.

1

ORIGIN

A name and a promise: important outputs should be stable enough to inspect.

2

ARCHITECTURE

Governance, logic, tone, divergence, and audit began appearing as separate layers.

3

MODULES

The same core idea was applied to law, weather, math, finance, hazards, and more.

4

CONTROL

The output path moved upstream: bind, check, replay, then release.

5

EVIDENCE

The record started carrying hashes, certificates, baselines, limits, and preserved artifacts.


The through-line

The point was never just to produce an answer. The point was to preserve enough structure around the answer to know what happened next.

EXHIBIT 01 / ORIGIN RECORD

The idea had an origin date.

A public post placed the beginning of Deterministic Intelligence on May 17, 2025.



WHAT THIS POST SHOWS

The image preserves a public origin story: a named idea, a date, and a claim that the work was about more than ordinary software determinism.


CAREFUL READING

This is chronology and public positioning. It is not an independent adjudication of invention, priority, or technical completeness.

EXHIBIT 02 / ARCHITECTURE

The architecture got a shape.

The public record started showing a runtime architecture instead of only describing an idea.



WHAT THIS POST SHOWS

A visual architecture record: inputs, logic, governance, tone, and a controlled output path are treated as separate but connected concerns.


CAREFUL READING

A diagram makes a system understandable. It does not, by itself, prove that every layer is fully implemented or universally effective.

EXHIBIT 03 / MODULE RECORD

The system became a set of modules.

By January, the public explanation was no longer one abstract engine. It was a family of named applications.



WHAT THIS POST SHOWS

The post presents a cross-domain module inventory and places auditability beside the product list rather than treating it as an afterthought.


CAREFUL READING

A module list is an inventory and a design statement. It is not proof that every named module has the same maturity or deployment status.

EXHIBIT 04 / SCROLL RECORD

Scrolls were explained as logic capsules.

The word scroll was used to mean a fixed, reusable unit of system logic.



WHAT THIS POST SHOWS

The post spells out the intended meaning: a scroll is not a decorative label. It is a bounded piece of logic, math, or engineering that can be reused and checked.


CAREFUL READING

This is a conceptual explanation of the architecture. It does not by itself establish code implementation, authorship priority, or legal protection.

EXHIBIT 05 / STACK RECORD

The stack learned to explain itself.

A seven-layer visual tried to show what happens between input and final output.



WHAT THIS POST SHOWS

The stack separates canonical input, governance, logic, tone, divergence handling, and artifact sealing. It is an attempt to make the control path readable.


CAREFUL READING

The visual communicates an architecture. It is not an independent performance test and does not establish that every output traverses every layer in every context.

EXHIBIT 06 / CROSS-DOMAIN RECORD

The same spine reached across domains.

The public posts connected legal, medical, financial, environmental, weather, and engineering examples.



WHAT THIS POST SHOWS

A cross-domain runtime is presented as one controlled substrate with domain-specific logic on top.

CAREFUL READING


Cross-domain architecture does not mean the domains have identical rules, evidence, risks, or validation standards.

Source: Social_Media_Posts_March_2026 / Cross-Domain_Runtime.jpg / commit e32663e91adb

EXHIBIT 07 / LEGAL WORKFLOW

The output path was written out.

BriefWise made the legal workflow visible as a sequence rather than a single answer box.



WHAT THIS POST SHOWS

The ladder illustrates a controlled legal path: define the situation, draft, review, test, gather authority, get feedback, and reach a final output.


CAREFUL READING

This is a workflow illustration. It is not itself a filing certification or a substitute for attorney review.

EXHIBIT 08 / CONTROL RECORD

Control moved upstream.

One of the clearest public explanations was also one of the simplest.



WHAT THIS POST SHOWS

The post states the design choice plainly: constrain behavior before the output is formed instead of trying to repair everything afterward.


CAREFUL READING

It is a design principle, not a guarantee that every downstream error disappears. Upstream controls still need evidence and testing.

EXHIBIT 09 / RUNTIME RECORD

A blocker became a next action.

The operating idea was not just to say no. It was to say what must happen next.



WHAT THIS POST SHOWS

The post describes a fail-closed workflow that preserves a blocker and routes the operator toward a concrete next step.


CAREFUL READING

A next-action display is not autonomous completion. It is a controlled handoff that still depends on the missing evidence, authority, or human decision.

EXHIBIT 10 / MATH RECORD

A fixed math case became a certificate.

MathWise used a narrow, replayable case instead of pretending that a local result was a general theorem.



WHAT THIS POST SHOWS

The post presents a fixed-case subset-sum result with a finite proof package and a clear target.


CAREFUL READING

The record is intentionally narrow. A fixed case is not the same thing as a general result, and a certificate is still open to independent review.

EXHIBIT 11 / MATH RECORD

The hard problem stayed scoped.

The Erdős work was presented as a reproducible certificate and research record, not as a magical general solution.



WHAT THIS POST SHOWS

A public post records a certificate-oriented mathematical investigation with a defined target and supporting package.


CAREFUL READING

This does not claim that the general Erdős problem was solved. The related Erdős 124 certificate prompt was a next ideal prompt, not a completed proof.

EXHIBIT 12 / REPLAY RECORD

Five artifacts matched.

The replay story became concrete when the same hashes appeared across different artifact layers.



WHAT THIS POST SHOWS

The post shows matching SHA-256 identities across CSV, NPZ, PNG, manifest, and final ZIP artifacts.


CAREFUL READING

This establishes artifact identity under the tested procedure. It does not prove that an underlying model is universally deterministic or that every analysis is correct.

EXHIBIT 13 / REPRODUCIBILITY RECORD

The three-run test was made concrete.

August added a compact visual for the same idea: identical final artifacts across three examined executions.



WHAT THIS POST SHOWS

The August README describes a three-execution Ulam artifact test in which the CSV, PNG, and complete ZIP matched byte-for-byte.

CAREFUL READING

The claim is deliberately limited to artifact-level reproducibility under the specified protocol. It is not a universal determinism claim.

EXHIBIT 14 / LEGAL EVIDENCE

A legal benchmark was recorded with a boundary.

The public post made a large result visible while labeling it as an internal, provisional audit.

Grounded DI LLC @Grounded... · 1h
 BriefWise DI² has now completed 501/501 provisional criteria across a series of blinded-rubric Harvey LAB benchmarks spanning materially different legal workflows. Results are public, internal audits, not official Harvey evaluations.

All-pass Rate only, ranked:

1. Grounded DI LLC / BriefWise DI² — 100.0% (501/501) provisional*
2. Claude Fable 5 (with fallback) — 14.2%
3. Claude Opus 4.8 (max) — 7.5%

[#DeterministicAI](#) [#AuditableAI](#) [#ai](#)

Grounded DI / BriefWise DI²
CERTIFICATE OF BENCHMARK AUDIT COMPLETION
 Harvey LAB Legal Work Product Series
 August 15, 2026

501 / 501 Criteria Locally Audited

This certificate represents completion of a multi-benchmark legal audit with product audit sequence completed a **Fallback 501** high score under the 501/501 criteria, and all 501/501. The sequence reached a provisional local continuity checkpoint of 501/501 criteria after completion of the 501/501 criteria. This is a certificate of benchmark audit completion and milestone recognition. It is not an official Harvey LAB or independent third-party audit, a high correctness guarantee, or a final Harvey evaluation.

Checkpoint Item	Value
Final score	501 / 501 (provisional)
Current benchmark	501 / 501 (provisional)
Current edition	501 / 501 (provisional)
Updated evidence	501 / 501 (provisional)
Required milestone	501 / 501 (provisional)
Finalized result	501 / 501 (provisional)
Package integrity	501 / 501 (provisional)
Finalized	501 / 501 (provisional)
Official Harvey evaluation	501 / 501 (provisional)

RUNTIME AND SOURCE BINDING

Field	Bound value
Runtime	501 / 501 (provisional)
Bound repository context	501 / 501 (provisional)
Test path	501 / 501 (provisional)
Bound auditability reference	501 / 501 (provisional)

Harvey LAB-AA: All-pass Rate
Share of tasks where every rubric criterion passes (Harvey's all-pass grading, no partial credit). Independently benchmarked by Artificial Analysis

Artificial Analysis

Claude Fable 5 (Adaptive Reasoning, Max Effort, Opus 4)
 14.2%

Click for more info

WHAT THIS POST SHOWS

The image preserves a public post reporting a 501/501 internal benchmark result and explicitly distinguishes it from an official external evaluation.

CAREFUL READING

The result is provisional and locally audited. It is not Harvey certification, independent legal correctness, or permission to use an output without professional review.

Source: -Social-Media-Posts-August-2026 / 1_BriefWise_DI_Provisional_Benchmark_Results_All-Pass_Rate_vs_External_Models.jpg / commit

Grounded DI LLC | public-source curation | source limits preserved

17

EXHIBIT 15 / AUDIT RECORD

Plausible data still needed an audit.

One August demonstration showed why synthetic data cannot be accepted merely because it looks coherent.



Grounded DI LLC

@Grounded_DI



Grounded DI OS did a one-shot audit of an external synthetic data repo for legal AI.

Although not necessarily representative, the demo reveals the challenges of using synthetic data without strong audit controls.

#AuditableAI #DeterministicAI #LawTech

By Grounded DI standards:

FilingReady: FAIL

Top reasons:

1. Material source conflicts

Facility totals and Pinnacle's MSA allocations differ across the workbooks and memoranda.

2. Arithmetic defects

Three HHI sheets contain negative "squared market-share" values. That makes the reported HHI results unreliable.

3. Unreconciled calculations

The underlying shares for "Other operators" are missing, so the HHI cannot be independently reproduced.

4. Conflicting methodology

The documents use different HHI classification thresholds, producing different labels for the same market.

5. Unresolved source authority

The record does not identify which facility inventory, market-share dataset, or HHI workpaper control.

Stated HHI	Stated ΔHHI	Data status
1,280	0	Supported by workbook components
1,150	0	Supported by workbook components
1,090	0	Supported by workbook components
980	0	Supported by workbook components
710	0	Conflicted
820	0	Supported by workbook components
1,210	0	Supported by workbook components
950	0	Conflicted
640	0	Conflicted

WHAT THIS POST SHOWS

The post records a one-repository audit that looks for conflicts, missing support, and mismatches inside synthetic legal data.

CAREFUL READING

It is one audited repository. It does not establish that every synthetic dataset has the same problems or that every flagged item is legally material.

Source: -Social-Media-Posts-August-2026 / Audit_of_One_Harveyai_Synthetic_Data_Repo.jpeg / commit 33429c250eda

EXHIBIT 16 / APPLICATION RECORD

The same idea reached interactive software.

The public record did not stay inside legal and research documents.




Grounded DI LLC
@Grounded_DI

The Tornado 🌀 on Mac — Visual physics powered by modified patent-pending weather logic.

This is one fun example of deterministic flexibility without reliance on machine learning.

#IndieGames #AI



WHAT THIS POST SHOWS

The Tornado is presented as a visual Mac application using modified weather logic in an interactive setting.

CAREFUL READING

A gameplay screenshot is a product demonstration, not proof that every component is machine-learning based, production-ready, or universally deterministic.

FINAL INTEGRITY CHECK

A strong record can still be a limited record.

The impressive part is the accumulation of specific, inspectable artifacts. The honest part is keeping each artifact at the level it actually proves.

THE RECORD SUPPORTS

- A sustained public design language for deterministic and auditable systems.
- A growing set of domain modules and workflow demonstrations.
- Specific artifact, replay, benchmark, and certificate records.
- A visible habit of naming missing inputs, unresolved states, and scope limits.
- A public progression from concept diagrams to working evidence packages.

THE RECORD DOES NOT AUTOMATICALLY PROVE

- Universal determinism of every underlying model or prompt.
- Independent validation of every public claim.
- A full solution to a general mathematical problem from a fixed-case certificate.
- Official endorsement by a benchmark owner or outside institution.
- Production readiness, legal sufficiency, or patent priority merely from a post.

That is not a weakness.

It is what makes the record usable. A result becomes more valuable when the boundary around it is visible.

SOURCE REGISTER / PART 1

Every exhibit has a source line.

Repository names and commit identities are included so the selection can be checked against the public record.

01 The idea had an origin date.

github.com/Grounded-DI/Social_Media_Posts_March_2026 | May_17_2025_The_"Big_Bang"_of_Deterministic_Intelligence.jpg | e32663e91adb

02 The architecture got a shape.

github.com/Grounded-DI/Social_Media_December_2025 | DIA_First_Deterministic_Runtime_Architecture.JPG | f101c65c1a9c

03 The system became a set of modules.

github.com/Grounded-DI/Social_Media_January_2026 | Grounded_DI_OS_Module_List_Plus_SHA-256_Auditability_Part_1.jpg | af697ecde624

04 Scrolls were explained as logic capsules.

github.com/Grounded-DI/Social_Media_Posts_February_2026 | DI_Principles_Scrolls_are_Logic_Capsules_for_Deterministic_Intelligence.jpg | a1aee00322b7

05 The stack learned to explain itself.

github.com/Grounded-DI/Social_Media_Posts_March_2026 | 7-Layer_Deterministic_Cognition_Stack.jpg | e32663e91adb

06 The same spine reached across domains.

github.com/Grounded-DI/Social_Media_Posts_March_2026 | Cross-Domain_Runtime.jpg | e32663e91adb

07 The output path was written out.

github.com/Grounded-DI/Social_Media_Posts_April_2026 | BriefWise_Workflow_Ladder_Visual.jpeg | 906f9ce5689f

08 Control moved upstream.

github.com/Grounded-DI/Social_Media_Posts_May_2026 | DI_Constrains_Behavior_Upstream.jpg | ca5d170d55af

SOURCE REGISTER / PART 2

Every exhibit has a source line.

Repository names and commit identities are included so the selection can be checked against the public record.

01 A blocker became a next action.

github.com/Grounded-DI/Social_Media_Posts_June_2026 | Grounded_DI_OS_It_Keeps_Converting_Each_Blocker_Into_Next_Action.jpeg | 0cfbc720984b

02 A fixed math case became a certificate.

github.com/Grounded-DI/Social_Media_Posts_May_2026 | MathWise_Subset-Sum_Covering_Theorem_(3_4_7_).jpeg | ca5d170d55af

03 The hard problem stayed scoped.

github.com/Grounded-DI/Social-Media-Posts-July-2026 | MathWise_Erdos_125_Reproducible_Certificate.jpeg | bde18a2314a4

04 Five artifacts matched.

github.com/Grounded-DI/Social-Media-Posts-July-2026 | 5-Hash-Reproducibility-Match.jpeg | bde18a2314a4

05 The three-run test was made concrete.

github.com/Grounded-DI/-Social-Media-Posts-August-2026 | FastPath_5.6_Instant_Three_Hash_Match_with_Deterministic_Visual.jpeg | 33429c250eda

06 A legal benchmark was recorded with a boundary.

github.com/Grounded-DI/-Social-Media-Posts-August-2026 | 1_BriefWise_DI_Provisional_Benchmark_Results_All-Pass_Rate_vs._External_Models.jpg | 33429c250eda

07 Plausible data still needed an audit.

github.com/Grounded-DI/-Social-Media-Posts-August-2026 | Audit_of_One_Harveyai_Synthetic_Data_Repo.jpeg | 33429c250eda

08 The same idea reached interactive software.

github.com/Grounded-DI/-Social-Media-Posts-August-2026 | The_Tornado_on_Mac_Live_Gameplay_Screenshot.jpeg | 33429c250eda

THE SIMPLE VERSION

Grounded DI kept making the invisible parts visible.

The posts did not all show the same kind of system. Some showed law. Some showed weather. Some showed mathematics. Some showed code, games, or hashes.

But the question stayed the same:

Can the output carry enough structure around it that another person can inspect what happened?

That is the evidence arc: from a promise about deterministic intelligence to a public collection of workflows, modules, audits, certificates, and replayable artifacts.

Precision does not guess. It traces.

Grounded DI LLC | Book 6 | The Evidence Arc | August 2026